

Pollux: Co-adaptive Cluster Scheduling for Goodput-Optimized Deep Learning

Seminar-Advanced Topics in AI for Networking

Aurick Qiao, Sang Keun Choe. et. al

Petuum, Inc. Carnegie Mellon University. UC Berkeley. MBZUAI.

15th USENIX Symposium & [Awarded best paper!](#)

August 16, 2022

Presented by Abdullah Amawi. University of Göttingen

Table of contents

- Introduction
 - Background
 - Problem Definition - Cluster challenges
 - Existing Solutions & Related works
- Pollux
 - Pollux Idea & Architecture
 - Pollux Scheduler Optimization
- Pollux System Modelling
- Evaluation
- Conclusion, Opinion & Future Work
- Questions & Discussion

Introduction

Background (Current Workflow)

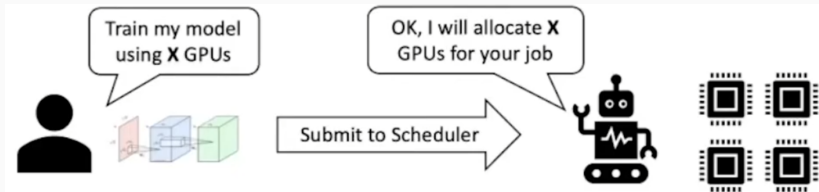


Figure 1: Shared-Cluster DL training workflow[1]

Considerations:

- **Cluster contention:** dynamic change based on usage
- **Job scalability:** which needs expert knowledge
- **Training parameters:** batch size & learning rate tuning

Background (Parallel Cluster DL)

- Distributed deep learning training requires a lot of optimizations for batch size and learning rate.

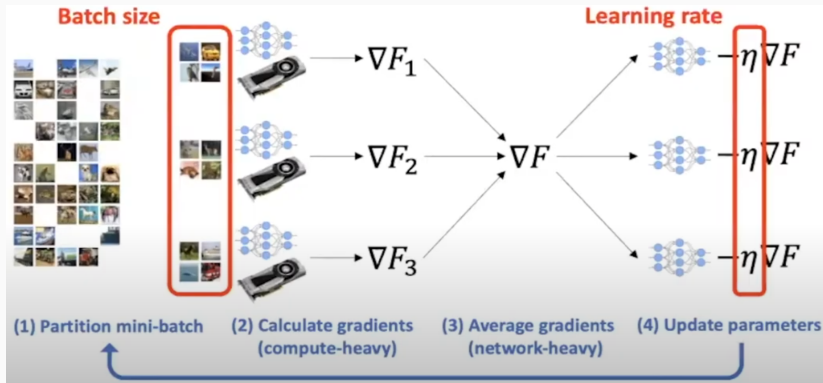


Figure 2: Distributed deep learning data parallelism[1]

Background(DL GPUs & Efficiency)

- Many GPUs doesn't translate to better throughput without optimizations; System **statistical efficiency** is still an issue!




Figure 3: Number of GPUs and system efficiency [1]

Background(DL Optimal batch size)

- **Batch size** relationship to training performance, system throughput and statistical efficiency **changes dynamically**.




Figure 4: Batch size & effects on different training stages[1]

Problem Definition - Cluster Scheduling Challenges

- Choosing good GPU allocation for DL depends on batch size, while batch size depends on the allocated GPUs, we have inter-dependency.
- GPU allocation is cluster-wide which relates to fairness and contention.
- Inter-dependant factors makes it very hard for users while submitting the jobs.

Existing Solutions & Related Works

- Existing DL schedulers [Tiresias & Optimus](#)[1] exist but with downfalls.
- [Adaptive DL schedulers](#): Optimus adapts the number of GPUs only, Tiresias doesn't adapt dynamically.
- [Adaptive batch size training](#): Many works exist[1] but only Pollux addresses realistic scenarios and [cluster scheduling](#).
- [Hyper_parameter tuning](#): Pollux aims to improve that with [dynamic adaptation](#) of parameters and inter-dependant factors.

Pollux.

Pollux revolves around these main ideas:

- **Dynamic tuning** of batch size and learning rate to utilize resources better
- Pollux introduces "**Goodput**", a new measure of DL training performance that combines **throughput w.r.t efficiency**.
- Pollux also considers **cluster-wide performance**, fairness and dynamically re-allocates resources to come up with better **Goodput**.

System Architecture

- Pollux [system Architecture](#) can be visualized by the following:

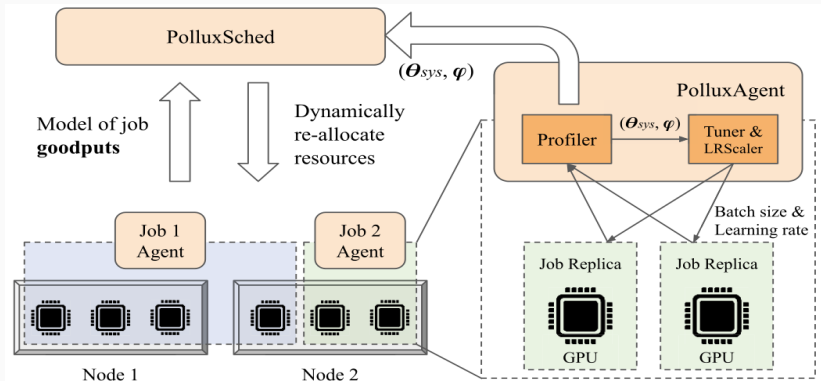


Figure 5: System architecture of Pollux[1].

Pollux Cluster Scheduler

- Pollux cluster scheduler illustration:

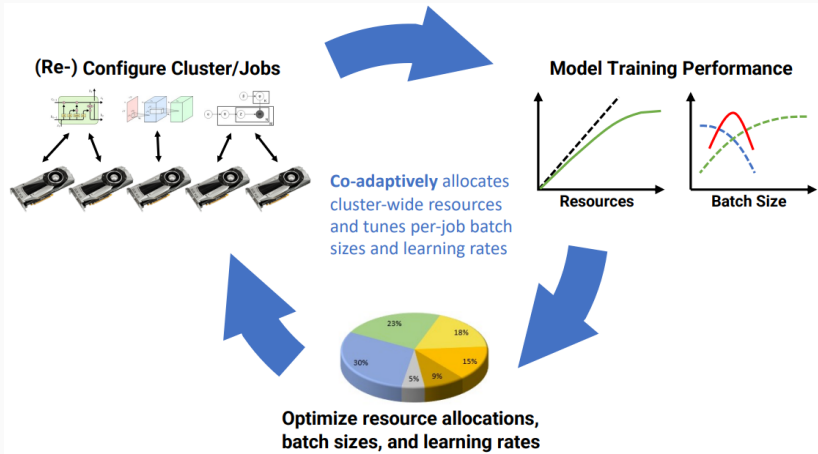


Figure 6: Pollux cluster scheduler illustration[1].

Pollux System Modelling

$$GOODPUT_t(a, m, s) = \underbrace{THROUGHPUT(a, m, s)}_{\text{System throughput (training examples / second)}} \times \underbrace{EFFICIENCY_t(M)}_{\text{Statistical efficiency (progress/training example)}}$$

- a : Allocation vector (Number of GPUs)
- m Per-GPU batch size
- s Gradient accumulation steps (to enable total batch sizes larger than GPU limit)
- M : Total batch size.
- a, m, s are automatically determined by Pollux.

Pollux System Throughput Modelling




Figure 7: System
throughput(Imagenet)[1]

- Models **throughput per GPU** count > **choose right number** of GPU and batch size.
- Models **gradients accumulation** > increase **batch size** beyond GPU **memory** limits.
- Models **GPU node allocation** > pack job's GPU onto fewer nodes to **minimize network overhead**.

Pollux Statistical Efficiency Modelling




Figure 8: Statistical efficiency(Imagenet)[1]

- Lower statistical efficiency for larger batch sizes > improves later in training.
- Accurately predicts statistical efficiency for different batch sizes > improves GOODPUT.
- Pollux predicts GOODPUT > even before running the job.

Evaluation & Results

- Pollux vs two SOTA DL schedulers(Optimus & Tiresias)
- Compared both in Makespan and average DL job completion time.
- Testbed
 - 16 AWS nodes w/ 64 GPUs (Nvidia T4, 4 each node)
 - 160 DL jobs submitted over 8 hours.
 - 48 CPUs, 192GB memory, and 900GB SSD.

Evaluation-Scheduling Vs Expert-tuned jobs

Policy	Avg job time	Makespan
Pollux (p = -1)	0.76h	16h
Optimus+Oracle+TunedJobs	1.5h	20h
Tiresias+TunedJobs	1.2h	24h

- 37-50% faster average training time in comparison to previous SOTA schedulers
- Note that Pollux is dynamic, compared to expert-configured jobs in competitors.

Evaluation - Statistical Efficiency

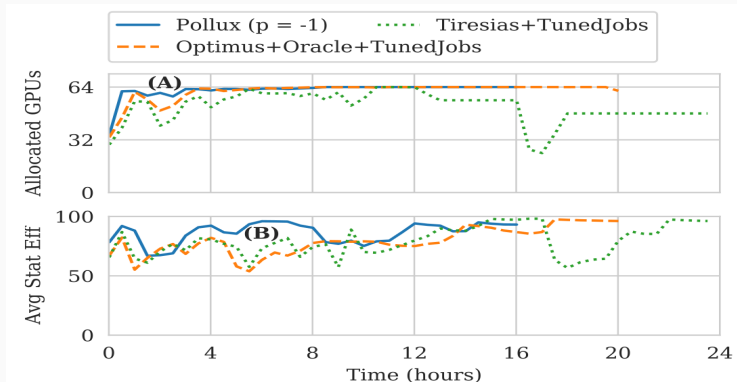


Figure 9: Pollux statistical cluster-wide efficiency [1]

- Reduces GPUs & batch size in high cluster contention(4-8)
- Accepts lower efficiency when we have lower cluster contention(8-12)

Conclusion, Opinion & Future Work

Conclusion.

- The paper presents Pollux, DL-aware cluster scheduler.
- Pollux co-optimizes both cluster-wide & per-job parameters for DL training.
- Pollux improves average DL training time in shared clusters by up to 50% even against very tuned baselines.
- Pollux co-adaptively allocates resources.

Opinion Of The Paper.

- On the positive side:
 - [Novel](#) idea that vastly improves DL jobs cluster scheduling.
 - [Clear presentation](#) of the problem and how Pollux solves it.
 - [First work](#) that considers GPU number, learning rate, and batch size dynamic job allocation in a cluster.
 - [Best paper](#) in a major conference.
- On the negative side:
 - [No future work](#) section; Some comments in the paper.
 - One note may be the [usage of Pytorch](#) and not [Tensorflow which is more production ready](#); But this is debatable.

Future Work

- [Tensorflow implementation](#) for being more production ready.
- Usage of different, more [capable GPUs](#).
- [Comparison](#) between own [private cloud & public cloud](#) such as the used AWS.
- Cloud [auto-scaling](#) system based on GOODPUT.
- [Full evaluation](#) on Pollux affects of different [hyper-parameter](#) algorithms.

Thank you!

Questions?

Additional resources

- **Goodput**: Measure for system throughput w.r.t statistical efficiency.

References



A. Qiao, S. K. Choe, S. J. Subramanya, W. Neiswanger, Q. Ho, H. Zhang, G. R. Ganger, and E. P. Xing.

Pollux: Co-adaptive cluster scheduling for goodput-optimized deep learning.

In 15th USENIX Symposium on Operating Systems Design and Implementation (OSDI 21), pages 1–18. USENIX Association, July 2021.